

Aria

A live AI teacher that draws while it speaks — and lets you interrupt it.

Product briefing · Use cases and market · US · July 2026

1. The one-line version

Every AI education product on the market **tells the student the answer**. Aria **teaches** — it talks, it draws on a board **while** it talks, and the student can cut in mid-sentence with “wait, why?” without it losing its place.

That sounds like a small distinction. It is the entire difference between a search engine and a teacher.

2. The story: what actually happens in a lesson

This is the workflow end to end, as a student experiences it.

Step 1 — The student says what they want to learn

They type it or say it: “Teach me photosynthesis.” No syllabus, no PDF to upload, no course to enrol in. Any topic, from a standing start.

Step 2 — Aria plans the whole lecture before it opens its mouth

Aria does not improvise sentence by sentence. It writes the **entire lecture as a sequence of teaching beats** first — the hook, the core idea, the worked example, the checkpoint question, the recap — so the lesson holds together from beginning to end and nothing repeats or contradicts itself later.

For each beat it decides **how that idea is best taught**: does this need a diagram drawn on the board? A real photograph? A question thrown back at the student?

This is the first thing nobody else does. A chatbot has no plan; it only responds. Aria has a lesson.

Step 3 — The board draws itself, in time with the voice

Aria begins speaking, and as it speaks **the board fills in underneath it**. The arrow appears at the exact moment it says “and this feeds into that.” A label is written the instant it is mentioned.

The board is not an image it generated and pasted up. It is being **drawn, stroke by stroke, in sync with the sentence being spoken** — the way a teacher at a chalkboard actually works.

Step 4 — The student interrupts

This is the heart of the product.

The student says “wait — why does the arrow go there?” Aria:

- **Stops** where it stands.
- **Circles the confusing thing on the board it already drew.** It does not wipe the board and start over.
- **Answers the actual question** — often by switching tactics. If the definition did not land, it tries an analogy.
- **Erases the annotation and picks the lesson up exactly where it left off.**

Compare that to what exists. A chatbot resets and replies with a wall of text. A video pauses. Neither can circle the thing you did not understand, because neither of them drew it.

Step 5 — Aria notices when the student is lost

It does not wait to be told. It watches the signals a human teacher watches: how long the student takes to answer a checkpoint, whether they ask the same thing twice, and — optionally — whether their attention has visibly drifted from the screen.

When it detects confusion it does not simply repeat itself louder. **It changes the explanation strategy** — swapping a formal definition for an analogy, turning a verbal explanation into a drawing, slowing the pace, or shortening the next chunk.

The attention sensing runs entirely on the student’s own device. No video ever leaves their browser.

Step 6 — The same lesson, rendered for a different mind

Here is where it compounds. Because Aria **composes** a lesson rather than rendering a picture, that one lesson can be delivered six different ways from a single source:

MODE	WHAT CHANGES
Standard	The board draws, the voice narrates, the student can interrupt at will.
Blind	The board is spoken aloud as it is drawn — “the arrow now runs from light to glucose” — with full voice control. The lesson is not reduced; it is re-rendered.
Dyslexia	Dyslexia-safe typography, re-paced narration, word-level highlighting synced to the voice.
Dysgraphia	The student never has to write in order to participate. Answering is by voice.
ADHD	Chunked into short beats, with the tutor actively re-engaging when attention drifts.
Autism	Predictable structure, sensory-load controls, literal-explanation mode.

A competitor whose slide is just an image cannot do this. Once a diagram has been flattened into a picture, the meaning is gone — you cannot speak it aloud to a blind student, because there is nothing left to speak. Aria never flattens it. That is not a feature we bolted on; it falls out of having built the thing correctly.

3. Why this is genuinely innovative

The market is crowded — but it is crowded in one corner

There are **2,800+ AI education startups operating in 2026**, an 18x increase since 2023, and they raised **\$4.2 billion in venture capital in 2025 alone**. AI tutoring attracted more venture funding than any other category in education AI.

But look at **what** they built. Content generation is by far the most crowded category — roughly **400 startups** doing essentially the same thing: generate a worksheet, a summary, a deck, a quiz.

Almost the entire market is racing to build **better answer machines**. Nobody is building a **teacher**.

The competitive map

PLAYER	WHAT IT DOES	WHAT IT CANNOT DO
Khanmigo	Socratic chat tutor over Khan Academy's curriculum. The strongest player in the category.	It is a chat window . It talks; it does not draw. Bound to Khan's existing content.
ChatGPT and general chatbots	Answers any question you can already articulate.	No lesson, no board, no pacing, no plan. You must already know what to ask.
MagicSchool, Diffit, ~400 others	Generate worksheets, summaries, reading-level-adapted text.	Teacher-facing content factories. They produce materials , not lessons.
Gamma, Tome, AI deck tools	Generate a finished slide deck and narrate over it.	A recording in a live costume. You can pause it — you cannot interrupt it.
Synthesia and AI video tools	Generate a talking-head video lecture.	Immutable. It cannot respond, adapt, or be asked a question.
Khan Academy's classic videos	Hand-drawn, dynamic, genuinely good teaching.	Pre-recorded . Fixed pace, one explanation, no interruption, no adaptation.
Aria	Plans a lecture, draws it live in sync with speech, is interruptible mid-sentence, adapts when the student is lost, and renders the same lesson for six different minds.	—

The evidence that drawing-while-speaking actually teaches better

This is not an aesthetic preference. It is one of the better-supported findings in multimedia learning research.

A **University of California, Santa Barbara** study had students watch the same explanation of the Doppler effect in two formats: **narration over a finished static diagram** — the PowerPoint model, and the model every AI deck tool has copied — versus **the same diagram being drawn by the instructor as they spoke**, the Khan Academy model.

The dynamic-drawing version produced measurably better retention and transfer. Showing the drawing being built engages the verbal and visual channels with complementary information, instead of forcing the student to split attention between a finished picture and a voice.

Here is the punchline: **every AI education tool on the market builds the losing format.** They generate a completed slide and talk over it. That is precisely the condition the research shows is worse.

Aria generates the winning format — and it is the only product that does so **live and interruptibly**, which even Khan Academy's videos cannot.

The innovation, stated plainly

Three things exist separately in the market today. **Nobody has combined them:**

- **Dynamic drawing synced to speech** — proven to teach better. Only Khan Academy does it, and only pre-recorded.
- **True conversational interruption** — chatbots have conversation but no lesson; videos have a lesson but no conversation.
- **One lesson, many minds** — accessibility as a **rendering** of the same lesson, not a separate and lesser product.

Aria is the intersection of the three. **That intersection is empty.**

4. Where it sells

Ranked by how quickly each converts to revenue.

1. Students with learning differences — dyslexia, ADHD, dysgraphia, autism, blindness

The wedge. Start here.

These families are already paying, already frustrated, and already underserved. Specialist tutors run **\$60–100 an hour**, often carry months of waitlist, and frequently are not available locally at all. And the specialist a family finally finds teaches **reading** — not chemistry, not algebra, not history.

Aria is a specialist tutor for **every subject**, available instantly, at a subscription price.

Why it is also the smartest wedge: it is the one market where our advantage is categorical rather than incremental. A chatbot cannot serve a blind student at all. We can. There is no incumbent to displace, because nobody is there.

2. Self-directed learners who want to genuinely understand something

The volume market.

The person teaching themselves a new subject — a career-switcher learning machine learning, a student who fell behind in calculus, a professional picking up finance — has two bad options today: read a wall of chatbot text, or watch a fixed-pace video that cannot hear their questions.

They do not want a summary. They want what a good teacher does: draw it, explain it, and let them ask.

This is where the market is large, where the “it feels like a real teacher” thesis either lands or does not, and where usage data comes from fastest.

3. Schools and universities — as a teaching assistant, not a replacement

The big contracts, and the slowest.

One teacher cannot deliver thirty simultaneous one-to-one explanations at thirty different paces. That is not a failing of teachers; it is arithmetic. Aria gives every student in the room a tutor that will re-explain the same concept a fifth time without impatience, in whatever way that particular student needs.

Institutions also badly want to serve students with learning differences and mostly cannot staff for it. We serve those students by default rather than as an add-on.

Reality check: institutional sales cycles run 9–18 months. Run this motion **behind** the consumer motion, funded by it.

4. Corporate training

Big budgets, fast decisions, low friction.

Corporate training today is a video and a quiz, completed by clicking “next” without watching. It is a ritual, not learning, and every L&D leader knows it. Aria turns it into a live, interruptible lesson where an employee can ask the question they are too embarrassed to ask in a room.

Enterprise budgets are large, procurement is quick, and there is none of the political friction of selling into schools.

5. Test prep and exam coaching

Later, but obvious.

The board **is** the product here — a maths or physics problem worked live, step by step, where the student can stop it at the exact line they lost the thread. Static explanations are terrible at this. Drawing is what it demands.

5. Honest status

Working today. Live lessons with a board that draws itself in sync with the voice; true mid-lesson interruption with annotate-and-resume; six delivery modes shipped (standard, blind, dyslexia, dysgraphia, ADHD, autism); voice interaction; on-device attention sensing; and lecture generation for any topic a student types in.

Not built yet. Full source-grounding and citations; accounts and saved progress; any efficacy evidence of our own.

This is a working product, not a mockup. But the gap between “working” and “sellable” is real, and closing it is what the next phase is for.

6. The ask

- **Two engineers for twelve weeks** — grounding and citations, accounts, and hardening.
- **A small efficacy study.** The research says dynamic drawing teaches better; we should be the ones who prove **our** version does. That evidence is the single most valuable sales asset we could own.
- **Ship the consumer motion first** — learning-difference families and self-directed learners. It earns revenue and usage data while institutional deals mature.

Sources

Grand View Research — AI Tutors Market Size, Share & Trends Report, 2026–2033

EduGenius — Education AI Startup Landscape 2026

Fiorella & Mayer (UC Santa Barbara) — dynamic drawing vs. static diagrams, *Journal of Educational Psychology*

CBE—Life Sciences Education — Effective Educational Videos: Principles and Guidelines

